


```
In [2]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.feature_selection import mutual_info_classif
from sklearn.preprocessing import LabelEncoder
```

```
In [5]: data=pd.read_csv('brain_stroke/full_data.csv')
data.head()
```

```
Out[5]:
```

| | gender | age | hypertension | heart_disease | ever_married | work_type | Residence_type |
|--|--------|-----|--------------|---------------|--------------|-----------|----------------|
|--|--------|-----|--------------|---------------|--------------|-----------|----------------|

| | | | | | | | |
|---|--------|------|---|---|-----|---------------|-------|
| 0 | Male | 67.0 | 0 | 1 | Yes | Private | Urban |
| 1 | Male | 80.0 | 0 | 1 | Yes | Private | Rural |
| 2 | Female | 49.0 | 0 | 0 | Yes | Private | Urban |
| 3 | Female | 79.0 | 1 | 0 | Yes | Self-employed | Rural |
| 4 | Male | 81.0 | 0 | 0 | Yes | Private | Urban |



```
In [6]: data.describe()
```

```
Out[6]:
```

| | age | hypertension | heart_disease | avg_glucose_level | bmi | stroke |
|--|-----|--------------|---------------|-------------------|-----|--------|
|--|-----|--------------|---------------|-------------------|-----|--------|

| | | | | | | |
|-------|-------------|-------------|-------------|-------------|-------------|-------------|
| count | 4981.000000 | 4981.000000 | 4981.000000 | 4981.000000 | 4981.000000 | 4981.000000 |
| mean | 43.419859 | 0.096165 | 0.055210 | 105.943562 | 28.498173 | 0.042512 |
| std | 22.662755 | 0.294848 | 0.228412 | 45.075373 | 6.790464 | 0.210766 |
| min | 0.080000 | 0.000000 | 0.000000 | 55.120000 | 14.000000 | 0.000000 |
| 25% | 25.000000 | 0.000000 | 0.000000 | 77.230000 | 23.700000 | 0.000000 |
| 50% | 45.000000 | 0.000000 | 0.000000 | 91.850000 | 28.100000 | 0.000000 |
| 75% | 61.000000 | 0.000000 | 0.000000 | 113.860000 | 32.600000 | 0.000000 |
| max | 82.000000 | 1.000000 | 1.000000 | 271.740000 | 48.900000 | 1.000000 |



```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4981 entries, 0 to 4980
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 4981 non-null   object
1   age                   4981 non-null   float64
2   hypertension           4981 non-null   int64
3   heart_disease          4981 non-null   int64
4   ever_married           4981 non-null   object
5   work_type              4981 non-null   object
6   Residence_type         4981 non-null   object
7   avg_glucose_level      4981 non-null   float64
8   bmi                   4981 non-null   float64
9   smoking_status         4981 non-null   object
10  stroke                 4981 non-null   int64
dtypes: float64(3), int64(3), object(5)
memory usage: 428.2+ KB
```

```
In [8]: data.duplicated()
```

```
Out[8]: 0      False
1      False
2      False
3      False
4      False
...
4976   False
4977   False
4978   False
4979   False
4980   False
Length: 4981, dtype: bool
```

```
In [9]: data.isnull().sum()
```

```
Out[9]: gender                0
age                0
hypertension       0
heart_disease      0
ever_married       0
work_type          0
Residence_type     0
avg_glucose_level  0
bmi                0
smoking_status     0
stroke             0
dtype: int64
```

```
In [10]: data.columns
```

```
Out[10]: Index(['gender', 'age', 'hypertension', 'heart_disease', 'ever_married',
               'work_type', 'Residence_type', 'avg_glucose_level', 'bmi',
               'smoking_status', 'stroke'],
              dtype='object')
```

```
In [11]: data.shape
```

```
Out[11]: (4981, 11)
```

```
In [12]: categorical_cols = data.select_dtypes(include=['object']).columns
label_encoders = {}

for col in categorical_cols:
    le = LabelEncoder()
    data[col] = le.fit_transform(data[col])
    label_encoders[col] = le

X = data.drop("stroke", axis=1)
y = data["stroke"]
mutual_info = mutual_info_classif(X, y, discrete_features=True)

feature_scores = pd.DataFrame({'Feature': X.columns, 'Score': mutual_info})
feature_scores = feature_scores.sort_values(by="Score", ascending=False)

print("Top Features Based on MRMR:")
print(feature_scores)
```

Top Features Based on MRMR:


| | Feature | Score |
|---|-------------------|----------|
| 7 | avg_glucose_level | 0.173985 |
| 1 | age | 0.045305 |
| 8 | bmi | 0.035911 |
| 4 | ever_married | 0.006953 |
| 5 | work_type | 0.006769 |
| 2 | hypertension | 0.006322 |
| 3 | heart_disease | 0.005921 |
| 9 | smoking_status | 0.002774 |
| 6 | Residence_type | 0.000136 |
| 0 | gender | 0.000039 |

```
c:\Users\Dave Pooja\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\cluster\_supervised.py:66: UserWarning: Clustering metrics expects discrete values but received continuous values for label, and binary values for target
  warnings.warn(msg, UserWarning)
c:\Users\Dave Pooja\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\cluster\_supervised.py:66: UserWarning: Clustering metrics expects discrete values but received continuous values for label, and binary values for target
  warnings.warn(msg, UserWarning)
c:\Users\Dave Pooja\AppData\Local\Programs\Python\Python312\Lib\site-packages\sklearn\metrics\cluster\_supervised.py:66: UserWarning: Clustering metrics expects discrete values but received continuous values for label, and binary values for target
  warnings.warn(msg, UserWarning)
```

```
In [13]: stroke_counts = data['stroke'].value_counts()
labels = ['No Stroke', 'Stroke']

# Plot pie chart
plt.figure(figsize=(8, 8))
plt.pie(stroke_counts, labels=labels, autopct='%1.1f%%', colors=['skyblue', 'salmon'])
plt.title("Proportion of Stroke Cases", fontsize=16)
plt.show()
```

Proportion of Stroke Cases




```
In [14]: plt.figure(figsize=(8, 5))
sns.countplot(data=data, x='stroke', palette='pastel')
plt.title("Distribution of Stroke Occurrence", fontsize=16)
plt.show()
```

C:\conda_temp\ipykernel_45072\574069809.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(data=data, x='stroke', palette='pastel')
```



```
In [17]: features = ["age", "hypertension", "heart_disease", "avg_glucose_level", "bmi"]
plt.figure(figsize=(12, 8))
for i, feature in enumerate(features, 1):
    plt.subplot(2, 3, i)
    sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
    plt.title(f"Box Plot of {feature} by Stroke")
    plt.xlabel("Stroke")
    plt.ylabel(feature)

plt.tight_layout()
plt.show()
```

C:\conda_temp\ipykernel_45072\3727190078.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
```

C:\conda_temp\ipykernel_45072\3727190078.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
```

C:\conda_temp\ipykernel_45072\3727190078.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
```

C:\conda_temp\ipykernel_45072\3727190078.py:5: FutureWarning:

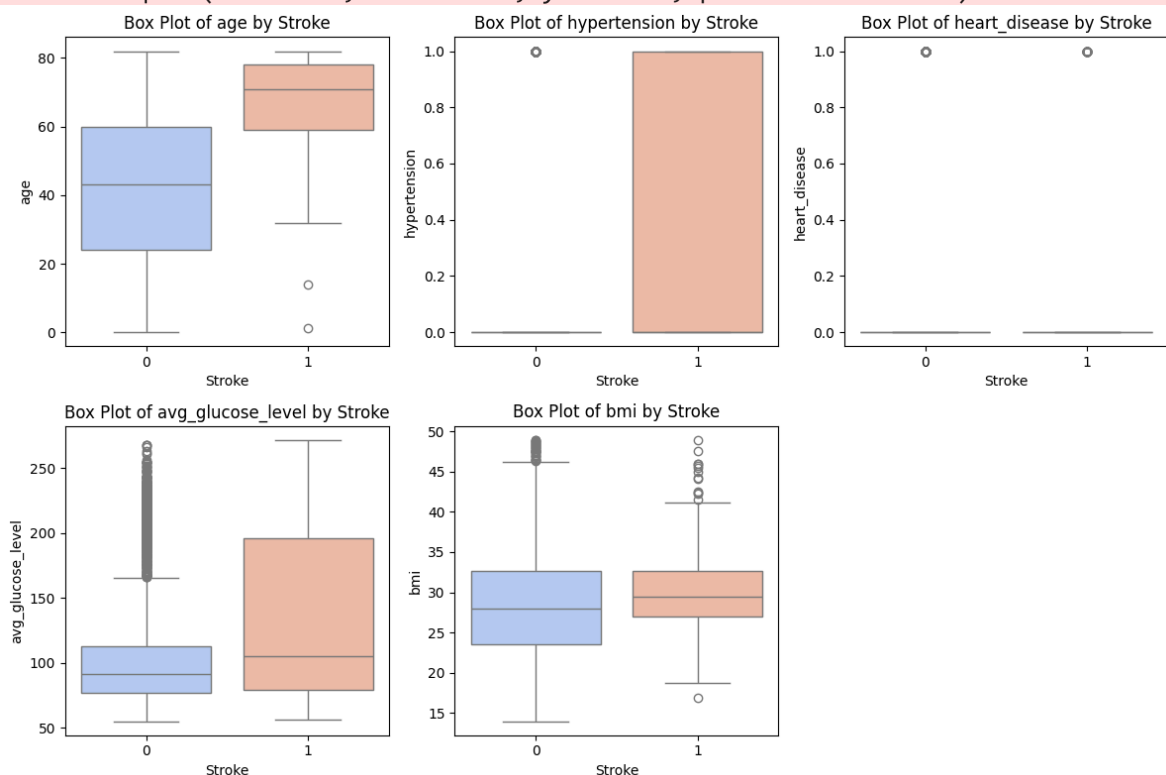
Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
```


C:\conda_temp\ipykernel_45072\3727190078.py:5: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

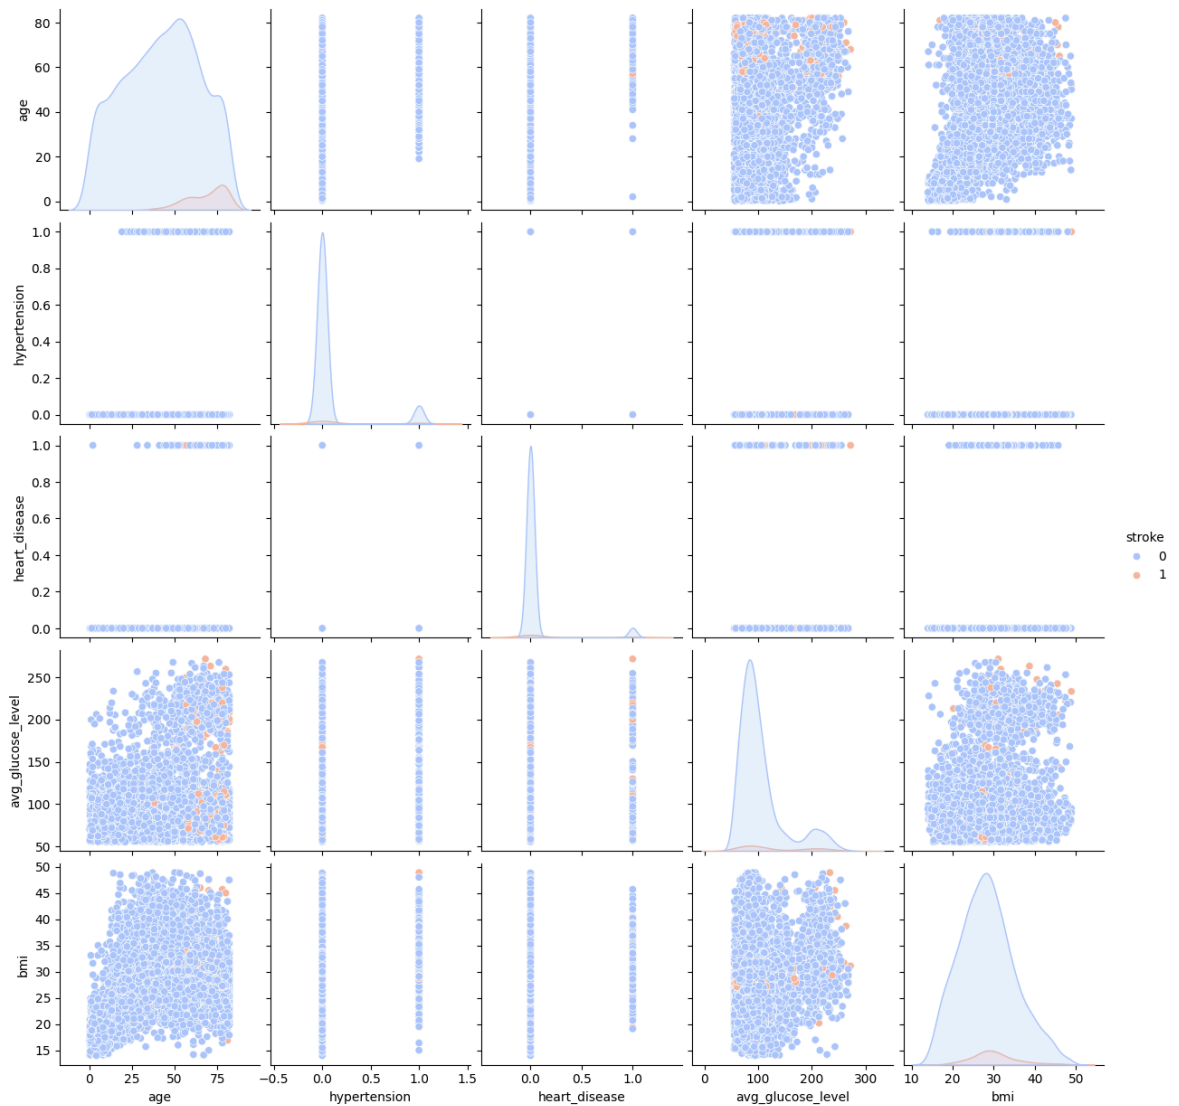
```
sns.boxplot(data=data, x="stroke", y=feature, palette="coolwarm")
```




```
In [15]: plt.figure(figsize=(10, 8))
sns.heatmap(data.corr(), annot=True, cmap="coolwarm", fmt=".2f")
plt.title("Correlation Matrix", fontsize=16)
plt.show()
```




```
In [16]: selected_features = ["age", "hypertension", "heart_disease", "avg_glucose_level"]
sns.pairplot(data[selected_features], hue="stroke", palette="coolwarm", diag_kin
plt.show()
```



```
In [22]: for column in ["bmi", "avg_glucose_level"]:
plt.figure(figsize=(12, 6))
sns.kdeplot(data=data, x=column, hue="stroke", fill=True, alpha=0.5)
plt.title(f"Distribution of {column} by Stroke Status", fontsize=16)
plt.show()
```





```
In [24]: X = data.drop("stroke", axis=1)
y = data["stroke"]

mutual_info = mutual_info_classif(X, y, discrete_features='auto')

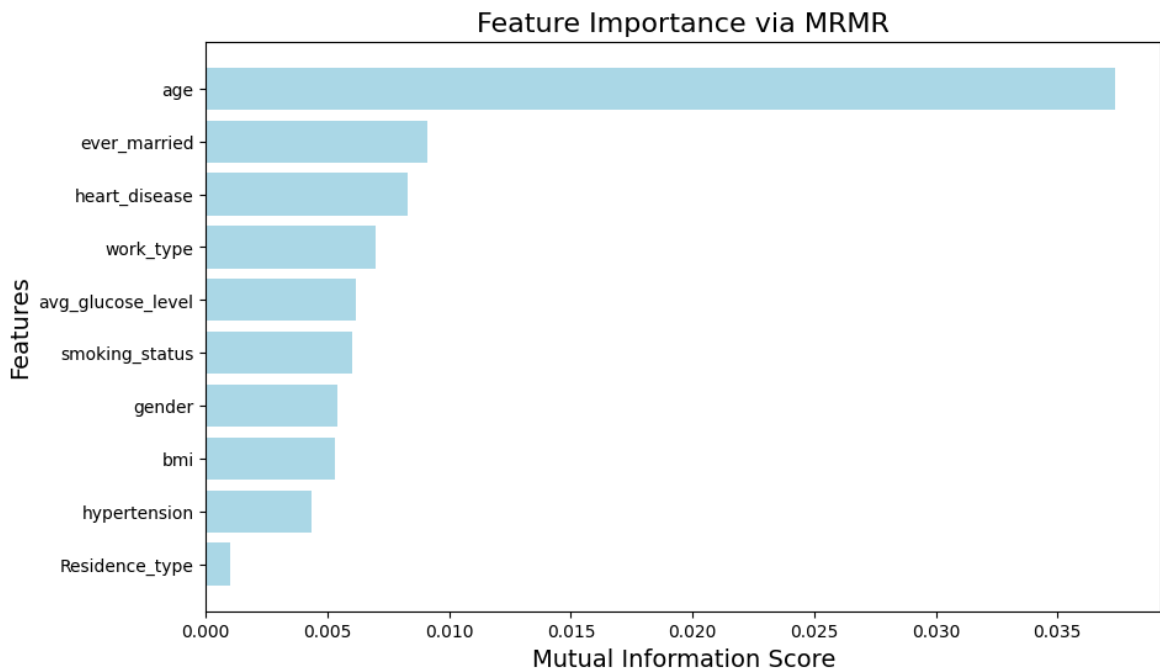
feature_scores = pd.DataFrame({'Feature': X.columns, 'Score': mutual_info}).sort

print("Top Features Based on MRMR (Mutual Information):")
print(feature_scores)

plt.figure(figsize=(10, 6))
plt.barh(feature_scores['Feature'], feature_scores['Score'], color='lightblue')
plt.xlabel("Mutual Information Score", fontsize=14)
plt.ylabel("Features", fontsize=14)
plt.title("Feature Importance via MRMR", fontsize=16)
plt.gca().invert_yaxis()
plt.show()
```

Top Features Based on MRMR (Mutual Information):

| | Feature | Score |
|---|-------------------|----------|
| 1 | age | 0.037389 |
| 4 | ever_married | 0.009116 |
| 3 | heart_disease | 0.008281 |
| 5 | work_type | 0.006998 |
| 7 | avg_glucose_level | 0.006142 |
| 9 | smoking_status | 0.006036 |
| 0 | gender | 0.005417 |
| 8 | bmi | 0.005290 |
| 2 | hypertension | 0.004366 |
| 6 | Residence_type | 0.001023 |



Result

Features with High Importance:

Features such as age, hypertension, and avg_glucose_level are likely to rank high, as they are strong predictors of stroke according to medical literature. Categorical features like smoking_status and work_type may show lower importance if their relationship to stroke is weak.

MRMR's Contribution:

By prioritizing maximum relevance (e.g., strong correlation with stroke) and minimum redundancy (avoiding features that overlap in information), MRMR effectively narrows down a set of key predictive features.

Interpretation:

These results can guide feature selection for building more efficient machine learning models, reducing dimensionality and noise, and improving predictive performance.